

Spark Practice

CONTENTS

1. Prerequisites.....	2
2. Task	3

PREREQUISITES

- Install Spark locally using one of the methods described [here](#) or in Docker.

I used Docker container to launch the Spark.

Containers [Give feedback](#)

Container CPU usage

1.31% / 800% (8 CPUs available)

Container memory usage

470.4MB / 7.57GB

Show charts

☰

☒ Only show running containers

<input type="checkbox"/>	Name	Image	Container ID	Port(s)	Actions
<input type="checkbox"/>	<div><div></div>container_spark</div>	jupyter/all-spark-notebook:latest	acbfc70346db	4040:4040 8888:8888	<div></div> <div></div> <div></div>

- Create a Spark ETL job to read data from a local storage. You can find the data in the Spark Practice—Dataset file on the page with the task description.

Downloaded the datasets to local disk and imported them by indicating the path in jupyter notebook.

The screenshot shows a Jupyter Notebook environment with a file explorer on the left and a terminal window on the right. The file explorer lists various files, including datasets, Spark_Task.ipynb, and test_run.ipynb. The terminal window displays the output of two code cells. Cell [2] shows the installation of pyspark, and Cell [93] shows the installation of geohash. The terminal output includes details about package requirements, downloads, and installation progress.

```
[2]: pip install pyspark

Requirement already satisfied: pyspark in /usr/local/spark/python (3.5.0)
Collecting py4j==0.10.9.7 (from pyspark)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl.metadata (1.5 kB)
  Downloading py4j-0.10.9.7-py2.py3-none-any.whl (280 kB)
    200.5/200.5 KB 846.3 kB/s eta 0:00:00a 0:00:01
Installing collected packages: py4j
Successfully installed py4j-0.10.9.7
Note: you may need to restart the kernel to use updated packages.

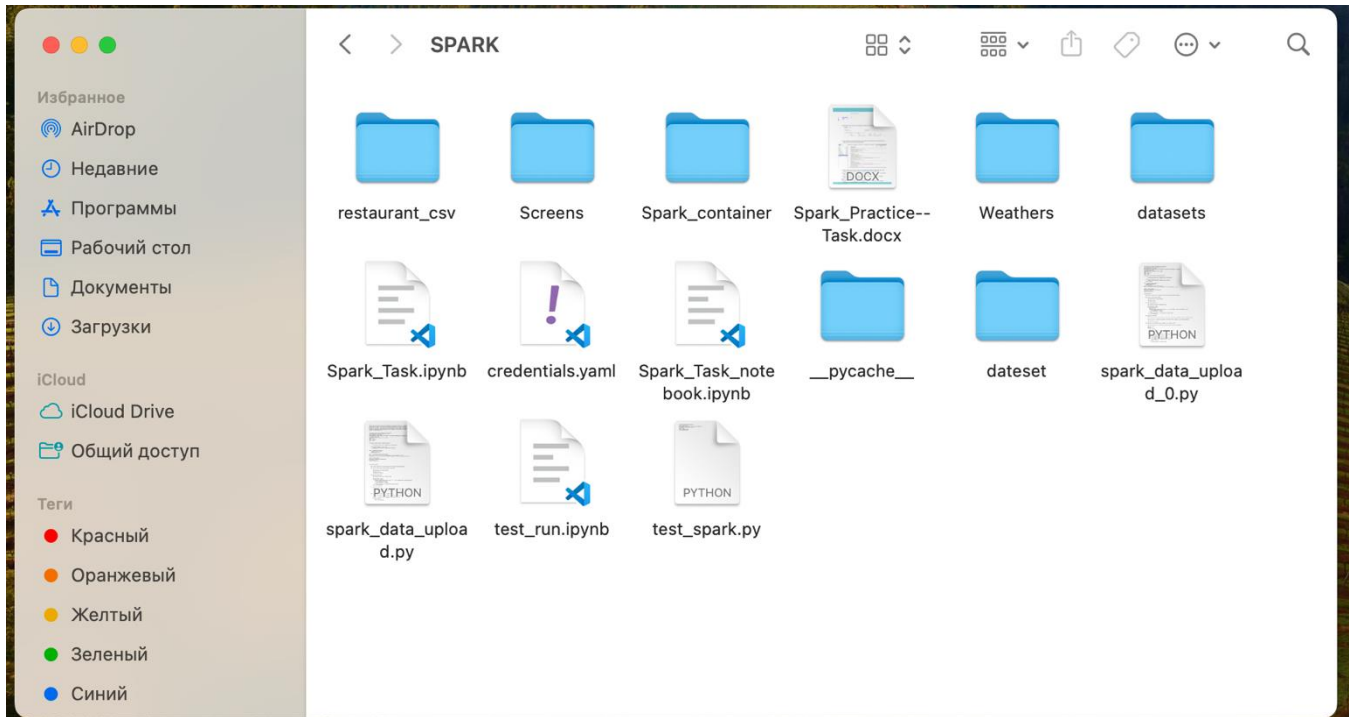
[93]: pip install geohash

Collecting geohash
  Downloading Geohash-1.0.tar.gz (15 kB)
  Preparing metadata (setup.py) ... done
Collecting docutils>=0.3 (from geohash)
  Downloading docutils-0.21.2-py3-none-any.whl.metadata (2.8 kB)
  Downloading docutils-0.21.2-py3-none-any.whl (587 kB)
    587.4/587.4 KB 348.2 kB/s eta 0:00:00a 0:00:01
Building wheels for collected packages: geohash
  Building wheel for geohash (setup.py) ... done
  Created wheel for geohash: filename=geohash-1.0-py3-none-any.whl size=15519 sha256=db796f422b17d187f1e4dc04d3bc61cedbb095fe2a3433a3e4daee41dbf
  Stored in directory: /home/jovyan/.cache/pip/wheels/b8/13/56/22a1fc3cd613735c1709eb859a780a99f287fc6fe0fc364b
Successfully built geohash
Installing collected packages: docutils, geohash
Successfully installed docutils-0.21.2 geohash-1.0
Note: you may need to restart the kernel to use updated packages.

[95]: pip install geohash2

Collecting geohash2
  Downloading geohash2-1.1.tar.gz (15 kB)
  Preparing metadata (setup.py) ... done
Requirement already satisfied: docutils>=0.3 in /opt/conda/lib/python3.11/site-packages (from geohash2) (0.21.2)
Building wheels for collected packages: geohash2
  Building wheel for geohash2 (setup.py) ... done
  Created wheel for geohash2: filename=geohash2-1.1-py3-none-any.whl size=15543 sha256=85f4e2bd238599c7b5f0ee4d147a3d01eb5709698da4d5fea80a4ff1f6de7
  Stored in directory: /home/jovyan/.cache/pip/wheels/f6/7c/c4/1b3c6feabebc3bf730dc86bbe7a173d581455df64c1f0623
Successfully built geohash2
Installing collected packages: geohash2
Successfully installed geohash2-1.1
Note: you may need to restart the kernel to use updated packages.

[63]: from pyspark.sql import SparkSession
```



TASK

- Check restaurant data for incorrect (null) values (latitude and longitude). For incorrect values, map latitude and longitude from the [OpenCage Geocoding API](#) in a job via the REST API.

```
import yaml
from tqdm import tqdm
from IPython.display import clear_output

[64]: spark = SparkSession.builder \
      .appName("CSVExample") \
      .getOrCreate()

[65]: path = '/home/jovyan/work/restaurant_csv'
      csv_files = [file for file in os.listdir(path) if file.endswith('.csv')]

[66]: df_restaurants = None
      for file in csv_files:
          #print(file)
          df_temp = spark.read.csv(path + '/' + file, header = True, inferSchema = True)
          if df_restaurants is None:
              df_restaurants = df_temp
          else:
              df_restaurants = df_restaurants.union(df_temp)

[67]: df_restaurants.show(5)

+-----+-----+-----+-----+-----+-----+
| id|franchise_id|franchise_name|restaurant_franchise_id|country|city|lat|lng|
+-----+-----+-----+-----+-----+-----+
|[257698037796]|37|Cafe Crepe|26468|IT|Milan|45.533|9.171|
|[25769803831]|56|The Waffle House|72238|FR|Paris|48.873|2.385|
|[85899345988]|69|Dragonfly Cafe|18952|NL|Amsterdam|52.392|4.911|
|[111669149758]|63|Cafe Paris|84488|NL|Amsterdam Zuidoo|52.31|4.942|
|[163288757268]|21|The Lazy Daisy|96638|US|Columbus|40.115|-83.015|
+-----+-----+-----+-----+-----+-----+
only showing top 5 rows

[68]: df_restaurants.count()

[68]: 1997

▼ Check Missing values

[69]: # Check for rows where latitude or longitude is null
      df_null_lat_lon = df_restaurants.filter(col("lat").isNull() | col("lng").isNull())
      # Show the rows with null latitude or longitude
      df_null_lat_lon.show()

+-----+-----+-----+-----+-----+-----+
| id|franchise_id|franchise_name|restaurant_franchise_id|country|city|lat|lng|
+-----+-----+-----+-----+-----+-----+
|[85899345928]|1|Savoria|18952|US|Dillon|NULL|NULL|
+-----+-----+-----+-----+-----+-----+
```

1. The process of reading dataset is done iteratively, due to several separate csv files. They were accumulated in an empty file. Dataset “Restaurant” contains 1997 rows, it was found 1 row with missing features latitude and longitude.
2. Filling the missing rows. API and Link of Service for Request is hidden in Credentials.yaml file.

Filling missing row

```
[71]: def load_config(file_path='credentials.yaml'):
      with open(file_path, 'r') as file:
          config = yaml.load(file, Loader=yaml.FullLoader)
      return config

[72]: config = load_config()
      url=config['url']
      api=config['api']

[73]: # Function to get latitude and longitude from OpenCage Geocoding API
      def get_lat_lon_from_address(address: str, api_key: str):
          base_url = url
          params = {
              'q': address,
              'key': api,
              'no_annotations': 1, # Skip unnecessary annotations
              'limit': 1          # Only get the top result
          }

          try:
              response = requests.get(base_url, params=params)
              data = response.json()

              # Check if the API returned results
              if data['results']:
                  lat = data['results'][0]['geometry']['lat']
                  lon = data['results'][0]['geometry']['lng']
                  return lat, lon
              else:
                  # Return None if no result found
                  return None, None
          except Exception as e:
              print(f"Error occurred while calling OpenCage API: {e}")
              return None, None
```

```
[74]: # Define UDF for getting lat/lon from address
def get_lat_lon_udf(address: str):
    # API Key for OpenCage Geocoding API
    API_KEY = api # Replace with your actual API key
    lat, lon = get_lat_lon_from_address(address, API_KEY)
    return (lat, lon)

# Register the UDF
new_lat_lon_udf_spark = udf(get_lat_lon_udf, returnType=StructType([
    StructField("lat", DoubleType(), True),
    StructField("lng", DoubleType(), True)
]))

[75]: #df_null_lat_lon

# Apply the UDF to fill in missing latitude/longitude values (assuming 'address' column exists)
df_updated = df_null_lat_lon.withColumn(
    "lat_lon", new_lat_lon_udf_spark(functions.col("city"))
)

# Split the 'lat_lon' column into 'latitude' and 'longitude' columns
df_updated = df_updated.withColumn("lat", df_updated["lat_lon"]["lat"]) \
    .withColumn("lng", df_updated["lat_lon"]["lng"]) \
    .drop("lat_lon")

# Show the updated DataFrame with corrected latitude and longitude
df_updated.show()
```

	id	franchise_id	franchise_name	restaurant_franchise_id	country	city	lat	lng
	85899345920	1	Savoria	18952	US	Dillon	34.4014089	-79.3864339

```
[14]: df_updated.createOrReplaceTempView("restaurant_data")

result = spark.sql('''SELECT * FROM restaurant_data
where city = 'Dillon' ''')

# Show the result of the query
result.show()
```

	id	franchise_id	franchise_name	restaurant_franchise_id	country	city	lat	lng
	85899345920	1	Savoria	18952	US	Dillon	34.4014089	-79.3864339

Once we fill the missing rows, we get modified and completed dataset with rows that had missing values.

- Then, I update the initial dataset with the actual data. I decided to continue my task in this way so that I don't need to transmit all the rows to be modified, and I minimize the risk of losing data or getting incorrect data.

```
[18]: '''
There are four steps that I will do
Alias the DataFrames to avoid ambiguity
Join the DataFrames on 'id'
Update lat and lon in df_restaurants with values from df_updated
Select columns from df_restaurants (using alias)
'''

df_restaurants_alias = df_restaurants.alias("restaurants")
df_updated_alias = df_updated.alias("updated")

df_joined = df_restaurants_alias.join(df_updated_alias, on='id', how='left')

df_updated_restaurants = df_joined.select(
    'id',

    df_restaurants_alias['franchise_id'],
    df_restaurants_alias['franchise_name'],
    df_restaurants_alias['restaurant_franchise_id'],
    df_restaurants_alias['country'],
    df_restaurants_alias['city'],
    # Use coalesce to get lat and lon from df_updated (if exists) or fallback to df_restaurants
    functions.coalesce(df_updated_alias['lat'], df_restaurants_alias['lat']).alias('lat'),
    functions.coalesce(df_updated_alias['lng'], df_restaurants_alias['lng']).alias('lng')
)

# Show the result (optional)
df_updated_restaurants.show()
```

	id	franchise_id	franchise_name	restaurant_franchise_id	country	city	lat	lng
	257698037796	37	Cafe Crepe	26468	IT	Milan	45.533	9.171
	25769803831	56	The Waffle House	72230	FR	Paris	48.873	2.305
	85899345988	69	Dragonfly Cafe	18952	NL	Amsterdam	52.392	4.911
	111669149758	63	Cafe Paris	84488	NL	Amsterdam Zuidoost	52.31	4.942
	163208757268	21	The Lazy Daisy	96638	US	Columbus	40.115	-83.015
	154618822662	7	Cafe Roma	41484	US	Tatum	33.382	-103.395
	163208757290	43	The Food House	96638	IT	Milan	45.474	9.224
	266287972361	10	The Golden Spoon	11263	US	Marina	36.684	-121.792
	171798691894	55	The Steak House	65939	GB	London	51.502	0.0
	197568495640	25	The Cozy Cafe	24784	US	Oskaloosa	41.324	-92.646
	163208757306	59	Azalea Cafe	96638	FR	Paris	48.871	2.294
	42	43	The Food House	47732	AT	Vienna	48.163	16.34
	23	24	The Fisherman's C...	47732	ES	Barcelona	41.396	2.163
	51539607572	21	The Lazy Daisy	6934	ES	Barcelona	41.387	2.174
	257698037775	16	The Spice Tree	26468	US	Morgantown	39.631	-79.956
	188978561036	13	The Firehouse	3642	US	Atlantic Beach	34.701	-76.747
	77309411367	40	Crimson Cafe	78190	AT	Vienna	48.213	16.357
	128849018902	23	The Hungry Pig	5679	ES	Barcelona	41.389	2.171
	223338299392	1	Savoria	36937	US	Washington	13.368	100.987

3.

- Generate a geohash by latitude and longitude using a geohash library like geohash-java. Your geohash should be four characters long and placed in an extra column.

Generating Geohash.

Geo Hash

```
[24]: def generate_geohash(lat, lon):
      return geohash.encode(lat, lon, precision=7)

geohash_udf = udf(generate_geohash, StringType())
df_with_geohash = df_updated_restaurants.withColumn("geohash", geohash_udf(col("lat"), col("lng")))
df_with_geohash.show(truncate=False)
```

id	franchise_id	franchise_name	restaurant_franchise_id	country	city	lat	lng	geohash
257698037796	37	Cafe Crepe	26468	IT	Milan	45.533	9.171	u0ne09n
25769803831	56	The Waffle House	72230	FR	Paris	48.873	2.305	u09wh3n
85899345988	69	Dragonfly Cafe	18952	NL	Amsterdam	52.392	4.911	u176pc8
111669149758	63	Cafe Paris	84488	NL	Amsterdam Zuidoost	52.31	4.942	u17986w
163208757268	21	The Lazy Daisy	96638	US	Columbus	40.115	−83.015	dphunyw
154618822662	7	Cafe Roma	41484	US	Tatum	33.382	−103.395	9tymzjn
163208757290	43	The Food House	96638	IT	Milan	45.474	9.224	u0nd9yk
266287972361	10	The Golden Spoon	11263	US	Marina	36.684	−121.792	9q92sw1
171798691894	55	The Steak House	65939	GB	London	51.502	0.0	gcpuzzx
197568495640	25	The Cozy Cafe	24784	US	Oskaloosa	41.324	−92.646	9zq55fc
163208757306	59	Azalea Cafe	96638	FR	Paris	48.871	2.294	u09wh0w
42	43	The Food House	47732	AT	Vienna	48.163	16.34	u2e9gzf
23	24	The Fisherman's Catch	47732	ES	Barcelona	41.396	2.163	sp3e3pz
51539607572	21	The Lazy Daisy	6934	ES	Barcelona	41.387	2.174	sp3e3qr
257698037775	16	The Spice Tree	26468	US	Morgantown	39.631	−79.956	dpp1kw9
188978561036	13	The Firehouse	3642	US	Atlantic Beach	34.701	−76.747	dq1mmt4
77309411367	40	Crimson Cafe	78190	AT	Vienna	48.213	16.357	u2edk0y
128849018902	23	The Hungry Pig	5679	ES	Barcelona	41.389	2.171	sp3e3qs
223338299392	1	Savoria	36937	US	Washington	13.368	100.987	w4ru418
103079215115	12	The Wooden Spoon	4340	US	Blythewood	34.214	−80.974	dnn6tkk

only showing top 20 rows

```
[25]: df_with_geohash.count()
```

```
[25]: 1997
```

- Left-join weather and restaurant data using the four-character geohash. Make sure to avoid data multiplication and keep your job idempotent.

Before the left join, I read the weather dataset iteratively by extracting each data and collected them into 1 spark dataframe

Import Weather Dataset

```
[26]: files_weather = [file for file in os.listdir('Weathers') if file.startswith('weather')]
```

```
[27]: df_w_all = None
for file_w in tqdm(files_weather):
    path_1 = 'Weathers/' + file_w
    path_2 = path_1 + '/' + os.listdir(path_1)[1]
    path_3 = path_2 + '/' + os.listdir(path_2)[1]
    days = [days for days in os.listdir(path_3) if days.startswith('day')]
    for day in days:
        path_4 = path_3 + '/' + day

        parquets = [parq for parq in os.listdir(path_4) if parq.endswith('.parquet')]
        for parquet in parquets:
            df_w_temp = spark.read.parquet(path_4 + '/' + parquet)

            if df_w_all is None:
                df_w_all = df_w_temp
            else:
                df_w_all = df_w_all.union(df_w_temp)
```

100%|██████████| 22/22 [00:08<00:00, 2.63it/s]

```
[28]: df_w_all.show(5)
```

lng	lat	avg_tmpr_f	avg_tmpr_c	wthr_date
-103.863	50.4005	64.1	17.8	2017-08-11
-103.799	50.4032	64.9	18.3	2017-08-11
-103.735	50.4058	64.2	17.9	2017-08-11
-103.671	50.4084	65.5	18.6	2017-08-11
-103.607	50.411	65.6	18.7	2017-08-11

only showing top 5 rows

```
[29]: df_w_all.count()
```

```
[29]: 112394743
```

Since the dataset is missing the Geohash, I generated geohash fro weather dataset.

Geo Hach for Weather

```
[30]: df_w_geohash_all = df_w_all.withColumn("geohash", geohash_udf(col("lat"), col("lng")))
```

```
[31]: # the result
df_w_geohash_all.show(truncate=False)
```

lng	lat	avg_tmpr_f	avg_tmpr_c	wthr_date	geohash
-103.863	50.4005	64.1	17.8	2017-08-11	c8ynsx1
-103.799	50.4032	64.9	18.3	2017-08-11	c8yntzx
-103.735	50.4058	64.2	17.9	2017-08-11	c8ynz2n
-103.671	50.4084	65.5	18.6	2017-08-11	c8yqbbt
-103.607	50.411	65.6	18.7	2017-08-11	c8yqf35
-103.543	50.4136	65.3	18.5	2017-08-11	c8yqgcd
-103.479	50.4162	65.2	18.4	2017-08-11	c8yqv3b
-103.415	50.4187	65.0	18.3	2017-08-11	c8yqydr
-103.351	50.4213	64.3	17.9	2017-08-11	c8ywb4y
-103.287	50.4238	63.9	17.7	2017-08-11	c8ywcck
-103.223	50.4263	63.6	17.6	2017-08-11	c8ywg5g
-103.159	50.4287	63.3	17.4	2017-08-11	c8ywws3
-103.095	50.4312	63.4	17.4	2017-08-11	c8ywyh8
-103.031	50.4336	63.4	17.4	2017-08-11	c8ywwmp
-102.966	50.436	63.0	17.2	2017-08-11	c8yybvww

```
[32]: df_grouped = df_w_geohash_all.groupBy("wthr_date", "geohash").agg(
    functions.avg("avg_tmpr_f").alias("avg_tmpr_f_avg"),
    functions.avg("avg_tmpr_c").alias("avg_tmpr_c_avg")
)

df_grouped.count()
```

```
[32]: 112394743
```

Note: Development and testing should be done locally in your IDE environment. Development and testing is proceeded locally in my machine(Jupyter lab)

- Store the enriched data (i.e., the joined data with all the fields from both datasets) in the local file system, preserving data partitioning in the parquet format.

The process of storing data was quiet challenging. Due to a large amount of data, it is better first decide in which way it is optimal in term of memory and time . Also, it depends on business task, if the business ask to provide specific group or segment of data, then better to do preprocessing and save only required data. In this task, we are asked to store all data, therefore, I splitted the whole data into 10 parts and saved separately as parquet file. **Parquet** is the most popular file format for Spark and many big data frameworks. It is a columnar storage format, meaning it stores data in columns rather than rows. Parquet supports efficient compression techniques, which reduce storage costs. It compresses better than row-based formats like CSV or JSON.

Join and saving table

```
[62]:
for idx, df1 in tqdm(enumerate(partitions)):
    clear_output(wait=True)
    print(f'Iteration {idx}. Data loading...')
    df1 = df1 \
        .withColumnRenamed("lat", "lat_1") \
        .withColumnRenamed("lng", "lng_1") \
        .withColumnRenamed("geohash", "geohash_1")

    df_joined = df1.join(
        df_with_geohash,          # The second DataFrame
        df1["geohash_1"] == df_with_geohash["geohash"], # The condition for the join
        "left"                   # The type of join (left join in this case)
    )

    df_joined.write.mode("overwrite").parquet(f'datasets/df_joined_{idx}')
    print(f'The parquet {idx} has succesfully recorded.')
```

```
Iteration 9. Data loading...
10it [25:19, 151.91s/it]
The parquet 9 has succesfully recorded.
```

■ / work / datasets /

Name	Last Modified
■ df_joined_0	11 hours ago
■ df_joined_1	11 hours ago
■ df_joined_2	11 hours ago
■ df_joined_3	10 hours ago
■ df_joined_4	10 hours ago
■ df_joined_5	10 hours ago
■ df_joined_6	10 hours ago
■ df_joined_7	10 hours ago
■ df_joined_8	10 hours ago
■ df_joined_9	10 hours ago
■ updated_restaurants	2 days ago
■ updated_restaurants.csv	2 days ago
■ updated_restaurants.parquet	2 days ago

■ / ... / datasets / df_joined_9 /

Name	Last Modified
📁 SUCCESS	10 hours ago
📁 part-00000-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00001-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00002-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00003-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00004-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00005-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00006-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00007-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00008-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago
📁 part-00009-cafae301-1676-416d-a4f5-4c8578983e42-c000.snappy.parquet	10 hours ago

You are expected to:

- Upload the source code and implement tests

Upload your fully documented homework with screenshots and comments in the task Readme MD file with the repo link. https://github.com/zhiyenbekov1222/Project_in_Spark.git

To completely finish the data, I crated a repository in GitHub and completed Readme MarkDown file. Also, I source code provided in **.py format and ready to run!**

The screenshot shows a GitHub repository named 'Project_in_Spark' by user 'zhiyenbekov1222'. The repository is private and has 13 commits. The main branch is selected. The file list includes README.md, spark_data_upload.py, test_spark.py, and unit_test.py. The README.md file is open, showing the title 'Spark Practice!' and a description of the project. The project is a part of the 'Data Engineering 2024 TechOrds' at Epam University. The project status is 'Completed'. The project intro/objective states that the purpose is to familiarize users with the basics of the open source distributed processing system for big data workloads. The methods used include Data Cleaning, API Integration, Geospatial Analysis, and Data Transformation. The right sidebar shows the repository's activity, including 0 stars, 1 watching, and 0 forks. It also lists suggested workflows for Django, Python package, and Python application.

github.com/zhiyenbekov1222/Project_in_Spark

zhiyenbekov1222 / Project_in_Spark

Code Issues Pull requests Actions Projects Security Insights Settings

Project_in_Spark Private

main 1 Branch 0 Tags

Go to file Add file Code

zhiyenbekov1222 Update README.md bf684d · 3 minutes ago 13 Commits

README.md Update README.md 3 minutes ago

spark_data_upload.py Add files via upload 5 days ago

test_spark.py Add files via upload 5 days ago

unit_test.py Add files via upload 3 minutes ago

README

Spark Practice!

Date completion: 02.12.2024

This project is a part of the [Data Engineering 2024 TechOrds](#) at Epam University.

Project Status: [Completed]

Project Intro/Objective

The purpose of this project is to be familiarize you with the basics of the open source distributed processing system for big data workloads. In addition to getting acquainted with the key components, architecture, and various applications of Spark, the project will discover the wealth of operations Spark offers, techniques about extract, transform, load (ETL), and the sets of APIs available in Spark. Apply the knowledge earned from the Spark lesson with the goal of building an understanding of how to improve the efficiency of Spark applications.

Methods Used

- Data Cleaning
- API Integration
- Geospatial Analysis
- Data Transformation

About Project_in_Spark

Readme Activity 0 stars 1 watching 0 forks

Releases No releases published [Create a new release](#)

Packages No packages published [Publish your first package](#)

Languages Python 100.0%

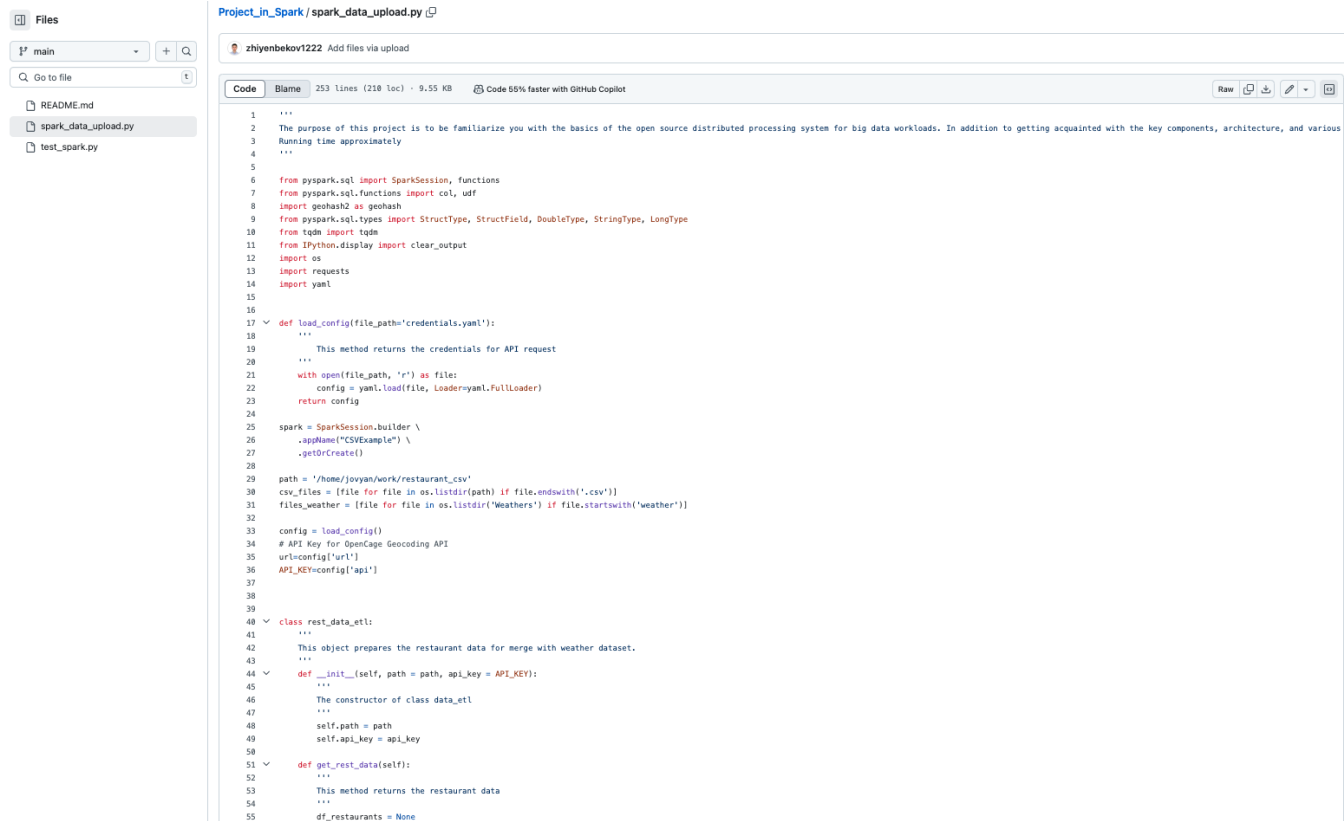
Suggested workflows Based on your tech stack

Django Configure Build and Test a Django Project

Python package Configure Create and test a Python package on multiple Python versions.

Python application Configure

LEARN. GROW. BE EXCELLENT.



```
1  ...
2  The purpose of this project is to be familiarize you with the basics of the open source distributed processing system for big data workloads. In addition to getting acquainted with the key components, architecture, and various
3  Running time approximately
4  ...
5
6  from pyspark.sql import SparkSession, functions
7  from pyspark.sql.functions import col, udf
8  import geohash2 as geohash
9  from pyspark.sql.types import StructType, StructField, DoubleType, StringType, LongType
10 from tqdm import tqdm
11 from IPython.display import clear_output
12 import os
13 import requests
14 import yaml
15
16
17 def load_config(file_path='credentials.yaml'):
18     ...
19     This method returns the credentials for API request
20     ...
21     with open(file_path, 'r') as file:
22         config = yaml.load(file, Loader=yaml.FullLoader)
23     return config
24
25 spark = SparkSession.builder \
26     .appName("CSExample") \
27     .getOrCreate()
28
29 path = '/home/jovyan/work/restaurant_csv'
30 csv_files = [file for file in os.listdir(path) if file.endswith('.csv')]
31 files_weather = [file for file in os.listdir('Weathers') if file.startswith('weather')]
32
33 config = load_config()
34 # API Key for OpenCage Geocoding API
35 url=config['url']
36 API_KEY=config['api']
37
38
39
40 class rest_data_etl:
41     ...
42     This object prepares the restaurant data for merge with weather dataset.
43     ...
44     def __init__(self, path = path, api_key = API_KEY):
45         ...
46         The constructor of class data_etl
47         ...
48         self.path = path
49         self.api_key = api_key
50
51     def get_rest_data(self):
52         ...
53         This method returns the restaurant data
54         ...
55         df_restaurants = None
```

The logs in testing the file. Totally, the running time of the process is roughly 27 minutes.

Below the screen...

LEARN. GROW. BE EXCELLENT.